


Homework 5

Benjamin Noland

1. (Cryer & Chan, Exercise 5.2)
2. (Cryer & Chan, Exercise 5.13)
 - a. The following R code produces a time series plot of the data:


```
temp_color = c(rev(brewer.pal(6, 'RdYlBu')), brewer.pal(6, 'RdYlBu'))  
  
plot(airpass, ylab = "Flights", type = "l")  
points(y = airpass, x = time(airpass), col = temp_color, pch = 15)
```



The plot seems to indicate a strong seasonal trend in the time series, with more flights in the warmer months than in the colder months. Moreover, the mean number of flights per year is steadily increasing over the whole time span of the series, as well as the variability in monthly flights within a given year – later years display considerably more variability than earlier years in the series.

- b. The following R code produces a time series plot of the log-transformed series:


```
plot(log(airpass), ylab = "log(Flights)", type = "l")  
points(y = log(airpass), x = time(airpass), col = temp_color, pch = 15)
```



The log transformation seems to have considerably dampened the yearly variability in the series. However, as in the untransformed case, the yearly means of the transformed data are steadily increasing, and the series displays the same seasonal trend.


- c. The following R code produces a time series plot of the differences of the log-transformed data:

```
diff_log <- diff(log(airpass))
plot(diff_log, ylab = "Difference of log(Flights)", type = "l")
points(y = diff_log, x = time(diff_log), col = temp_color, pch = 15)
```



Next, we have a time series plot of the (fractional) relative differences:

```
rel_change <- (airpass / zlag(airpass)) - 1
plot(rel_change, ylab = "Relative change", type = "l")
points(y = rel_change, x = time(rel_change), col = temp_color, pch = 15)
```



The two plots seem to display the same, relatively stable trend. Assuming the true model is of the form

$$Y_t = (1 + X_t)Y_{t-1},$$

with $|X_t|$ small for every t , then

$$\nabla \log Y_t = \log \left(\frac{Y_t}{Y_{t-1}} \right) = \log(1 + X_t) \approx X_t$$

for every t . Moreover,


$$\frac{Y_t - Y_{t-1}}{Y_{t-1}} = \frac{Y_t}{Y_{t-1}} - 1 = (1 + X_t) - 1 = X_t.$$

Thus, under this model, we would expect the two plots to look about the same.

3. (Cryer & Chan, Exercise 5.14)

- a. The following R code generates a plot of log-likelihood against λ .

```
BoxCox.ar(larain)
```




The vertical dashed lines denote a 95% confidence interval for λ . In particular, from the plot we see that the value $\lambda = 0.25$ lies within this confidence interval. For the sake of potential interpretability, we will take this to be the value of λ for the power transformation.

- b. The following R code produces two plots: the left plot is a normal Q-Q plot of the original (untransformed) data, and the right plot is one of the transformed data.

```
lambda <- 0.25
transformed <- (larain^lambda - 1) / lambda

old_par <- par(mfrow = c(1, 2))
qqnorm(larain, main = "Original data")
qqline(larain)


qqnorm(transformed, main = "Transformed data")
qqline(transformed)
par(old_par)
```




The normal Q-Q plot of the transformed data adheres very closely to the theoretical Q-Q plot, indicated by the solid black line. Thus, in contrast with the original data, the transformed data appear to be normal.

- c. The following R code produces time series plots of both the original (untransformed) data and the transformed data:

```
plot(larain, type = "o", main = "Original data",
     xlab = "Year", ylab = "Inches")
plot(transformed, type = "o", main = "Transformed data",
     xlab = "Year", ylab = "Inches")
```






Notice that the correlations between the data points appear to be similar in both plots.

- d. The following R code plots Y_t vs. Y_{t-1} for both the original (untransformed) data and the transformed data.

```
old_par <- par(mfrow = c(1, 2))
plot(y = larin, x = zlag(larin), main = "Original data",
     xlab = "Previous year inches", ylab = "Inches")
plot(y = transformed, x = zlag(transformed), main = "Transformed data",
     xlab = "Previous year inches", ylab = "Inches")
par(old_par)
```



Both plots seem to indicate a lack of correlation between Y_t and Y_{t-1} . Moreover, we should not expect the transformation to change the dependence between data points in this series. This is a consequence of the following fact: if Y_1, \dots, Y_n are random variables and g is a (measurable) function, then $g(Y_1), \dots, g(Y_n)$ are random variables, and they are mutually independent if and only if Y_1, \dots, Y_n are mutually independent. In this case, Y_1, \dots, Y_n are elements of the time series, and $g(x) = (x^\lambda - 1)/\lambda$, where $\lambda = 0.25$.

5. The following R function generates n observations from the MA(3) process in question using Gaussian noise:

```
set.seed(1) # Set a random seed for reproducibility

sample_process <- function(n) {
  wn <- ts(rnorm(n), start = -2, end = n)
  process <- wn - 0.7 * zlag(wn, 1) - 0.5 * zlag(wn, 2) - 0.6 * zlag(wn, 3)
  process <- window(process, start = 1) # Scrap times < 1
  process
}
```

The following is a time series plot of the the resulting simulated data for $n = 500$:

```
plot(sample_process(500), type = "l", main = "Simulated process",
      ylab = "Value", lwd = 0.5)
```

